

Per-paper performance, sorted by cross-model mean
(each bar = one paper's experiment-macro score for one model)

